

# Part 1

## AI/ML Project Lifecycle

A comprehensive guide to building and deploying machine learning solutions

### Problem Definition

#### Hypothetical AI Problem

Predicting Student Dropout Rates

Building a predictive model to identify students at risk of dropping out, enabling early intervention and support.

### Objectives

- Achieve 85%+ prediction accuracy for at-risk students
- Enable proactive intervention strategies for educational institutions
- Reduce overall dropout rates by identifying early warning signs

### Stakeholders

- Educational Institutions: Schools and universities implementing support programs
- Students & Families: Direct beneficiaries of timely interventions

### Key Performance Indicator (KPI)

Dropout Rate Reduction

Measure: Percentage decrease in dropout rates within 12 months of model deployment compared to historical baseline.

### Data Collection & Preprocessing

### Data Sources

- Student Information System (SIS): Academic records, attendance, grades, demographics

- Learning Management System (LMS): Engagement metrics, assignment completion, online activity patterns

## Potential Data Bias

Socioeconomic Bias: Historical data may overrepresent certain socioeconomic groups, leading to biased predictions. Students from underrepresented backgrounds might be incorrectly flagged or missed due to systemic data collection gaps.

## Preprocessing Steps

- Handling Missing Data: Impute missing values using mean/median for numerical features and mode for categorical features, or use advanced methods like KNN imputation
- Feature Normalization: Apply standardization (z-score) or min-max scaling to numerical features to ensure all features contribute equally to model training
- Categorical Encoding: Convert categorical variables (gender, major, ethnicity) using one-hot encoding or label encoding for model compatibility

## Model Development

### Model Choice & Justification

#### Random Forest Classifier

Justification: Random Forest is well-suited for this problem because it:

- Handles both numerical and categorical features effectively

- Provides feature importance rankings to identify key dropout predictors
- Resistant to overfitting through ensemble averaging
- Performs well with imbalanced datasets (common in dropout scenarios)

## Data Splitting Strategy

Training Set (70%): Used to train the model on historical patterns

Validation Set (15%): Used for hyperparameter tuning and model selection without touching test data

Test Set (15%): Final evaluation on unseen data to assess real-world performance

*Use stratified sampling to maintain class distribution across all sets, especially important for imbalanced dropout data.*

## Hyperparameter Tuning

- n\_estimators (Number of Trees): More trees generally improve performance but increase computation time. Tune to find the sweet spot between accuracy and efficiency (typical range: 100-500).
- max\_depth (Tree Depth): Controls model complexity and prevents overfitting. Deeper trees capture more patterns but risk memorizing noise. Optimize to balance training and validation performance (typical range: 5-20).

## Evaluation & Deployment

## Evaluation Metrics

- Recall (Sensitivity): Measures the proportion of actual at-risk students correctly identified. Critical in this context because missing a student who drops out has significant consequences. High recall ensures we catch most at-risk students.
- F1-Score: Harmonic mean of precision and recall, providing a balanced view. Important because we need to balance catching at-risk students (recall) while not overwhelming support services with false positives (precision).

## Concept Drift

What is Concept Drift?

Concept drift occurs when the statistical properties of the target variable change over time, causing model performance to degrade. In student dropout prediction, this could happen due to:

- Changes in curriculum or teaching methods
- Economic shifts affecting student circumstances
- New support programs altering dropout patterns

Monitoring Strategy:

- Track model accuracy, recall, and F1-score monthly
- Compare predictions against actual outcomes continuously
- Set up alerts when metrics drop below defined thresholds
- Retrain model quarterly or when significant drift is detected

## **Technical Deployment Challenge**

### **Scalability & Real-Time Processing**

**Challenge:** The system must process predictions for thousands of students across multiple institutions simultaneously, with near real-time updates as new data arrives.

### **Solution Approach:**

- Implement batch processing for routine predictions (daily/weekly)
- Use containerization (Docker/Kubernetes) for horizontal scaling
- Employ caching strategies for frequently accessed predictions
- Design asynchronous architecture to handle peak loads during enrollment periods